

截图 1:

```
(client) measure@LAPTOP-CM1:~$ curl -s http://localhost:30000/v1/chat/completions \
-H "Content-Type: application/json" \
-d '{
  "model": "/home/measure/models/Qwen3-0.6B",
  "messages": [{"role": "user", "content": "用一句话介绍一下你自己。"}],
  "max_tokens": 256,
  "chat_template_kwargs": {"enable_thinking": false}
}' | python -m json.tool

{
  "id": "2450195ab6cb44a8bdfc6d9c86dd56de",
  "object": "chat.completion",
  "created": 1790410430,
  "model": "/home/measure/models/Qwen3-0.6B",
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "\u6211\u662fAI\u52a9\u624b\u7684\u529f\u80fd\u662f\u5e2e\u52a9\u7528\u623f\u89e3\u7b54\u95ee\u9898\u6216\u63d0\u4f9b\u652f\u6301\u3002",
        "reasoning_content": null,
        "tool_calls": null
      },
      "logprobs": null,
      "finish_reason": "stop",
      "matched_stop": 151645
    }
  ],
  "usage": {
    "prompt_tokens": 17,
    "total_tokens": 30,
    "completion_tokens": 13,
    "prompt_tokens_details": null,
    "reasoning_tokens": 0
  },
  "metadata": {
    "weight_version": "default"
  }
}
```

发送一次请求推理 POST /v1/chat/completions

用户：用一句话介绍一下你自己。

模型回答内容：

\u6211\u662fAI\u52a9\u624b\u7684\u529f\u80fd\u662f\u5e2e\u52a9\u7528\u623f\u89e3\u7b54\u95ee\u9898\u6216\u63d0\u4f9b\u652f\u6301\u3002（我是 AI 助手，可以帮你解答问题或提供支持。）

截图 2 在下一页(●' ∪' ●)

截图 2:

```
(client) measure@LAPTOP-CMJ:~/trace/Mooncake/FAST25-release$ python "/mnt/d/AAACMJ/study/fzu/节假日作业训练营/第一关/scripts/03_mooncake_bench.py" \
--trace ~/trace/Mooncake/FAST25-release/xxiv-trace/mooncake_trace.jsonl \
--n 20 --seed 42 --rate 1.0 \
--min-input 890 --max-input 2048 --max-output 256 \
--out run1.csv --tag run1
[info] 读取 ~/home/measure/trace/Mooncake/FAST25-release/xxiv-trace/mooncake_trace.jsonl
[info] 过滤后剩 980 条 (input 890-2048, output<= 256)
[transformers] PyTorch was not found. Models won't be available and only tokenizers, configuration and file/data utilities can be used.
[info] 已加载 tokenizer: Qwen/Qwen3-0.6B (token 数将精确对齐)
[info] 热身完成
[info] 20 条请求已按泊松过程发出, 等待完成.....

Mooncake Trace -> SGLang 回放 tag=run1 model=Qwen3-0.6B
trace=~/home/measure/trace/Mooncake/FAST25-release/xxiv-trace/mooncake_trace.jsonl
n=20 rate=1.0 req/s seed=42 input=[890,2048] output<= 256

=====
# | 计划ln | 实际ln | 计划out | 实际out | 状态 | TTFT(ms) | 总延迟(ms) | TPOT(ms) | 到达(s)
-----
1 | 956 | 969 | 189 | 189 | success | 43.3 | 1064.2 | 5.4 | 0.00
2 | 979 | 992 | 3 | 3 | success | 69.6 | 78.0 | 4.2 | 0.29
3 | 1179 | 1193 | 27 | 27 | success | 91.0 | 230.1 | 5.7 | 0.73
4 | 1943 | 1958 | 21 | 21 | success | 136.3 | 262.3 | 6.3 | 3.46
5 | 1395 | 1409 | 86 | 86 | success | 62.8 | 583.8 | 6.1 | 5.62
6 | 894 | 907 | 63 | 63 | success | 37.9 | 485.0 | 5.9 | 6.29
7 | 1927 | 1942 | 19 | 19 | success | 49.3 | 160.0 | 6.6 | 12.12
8 | 1454 | 1468 | 116 | 116 | success | 65.4 | 968.5 | 7.9 | 13.55
9 | 1972 | 1987 | 62 | 62 | success | 119.5 | 573.5 | 7.6 | 14.20
10 | 1348 | 1362 | 64 | 64 | success | 89.4 | 454.6 | 5.8 | 14.39
11 | 897 | 910 | 216 | 216 | success | 24.4 | 1056.4 | 4.8 | 15.38
12 | 1392 | 1406 | 146 | 146 | success | 40.1 | 711.3 | 4.6 | 16.71
13 | 910 | 923 | 26 | 26 | success | 27.7 | 148.1 | 4.8 | 18.83
14 | 2036 | 2051 | 22 | 22 | success | 61.5 | 160.8 | 4.7 | 19.16
15 | 892 | 905 | 211 | 211 | success | 23.1 | 1012.4 | 4.7 | 21.23
16 | 891 | 904 | 220 | 220 | success | 29.5 | 1048.3 | 4.7 | 21.67
17 | 931 | 944 | 21 | 21 | success | 35.4 | 126.5 | 4.6 | 23.11
18 | 890 | 903 | 107 | 107 | success | 23.2 | 486.7 | 4.4 | 25.00
19 | 1213 | 1227 | 60 | 60 | success | 40.8 | 330.0 | 4.9 | 26.25
20 | 1179 | 1193 | 236 | 236 | success | 86.9 | 1183.0 | 4.7 | 29.11
-----

汇总: 成功 20/20 | 平均 TTFT 57.4ms | P50 TTFT 46.3ms | P95 TTFT 136.3ms
平均延迟 553.0ms | 总输入 tokens 25,553 | 总输出 tokens 1,915
回放总时长 30.30s | 输出吞吐 63.2 tok/s | 请求吞吐 0.66 req/s
input 偏差: 均值 +13.8 token (chat template 开销)
output 偏差: 均值 +0.0 token (ignore_eos 下应接近 0)
=====
```

从 Mooncake FAST’25 trace(mooncake_trace.jsonl)中采样 20 条请求记录
记录每个请求的 input tokens、output tokens、status、TTFT、latency
程序 03_mooncake_bench.py 由 Claude Code 生成
包含:

- 1. 读 trace + 过滤 + 采样
- 2. 构造 synthetic prompt
- 3. 泊松到达时刻
- 4. 发送 + 计时